


Majority Voting & Ensemble Learning

1. Example

- KNN (k=3: 3 nearest neighbors)



3 votes [1,3,3] for final class decision: **class label = 3.**

1. Example

- Random Forest

	DT ₁	DT ₂	DT ₁₀₀
X	l ₁	l ₁	l ₁₀₀

100 votes [l₁, l₂, ..., l₁₀₀] for final class decision

- One decision tree (DT₁) is described as follows

C ₁	C ₂	C ₁₀
0.15	0.8	0.9

Probabilities of being classified to each class

1. Example

- Combination/Fusion of different methods

SVM	KNN	Decision Tree	Linear Regression	Logistic Regression
l1	l2	l3	l4	l5

5 votes [l1, l2, l3, l4, l5] for final class decision

2. Two-class Voting

- Decision Tree (DT)

	Correct	Incorrect
$p=$	0.5	0.5
$p=$	0.6	0.4

Bad Tree

Good Tree

Probabilities of being classified to correct/incorrect class

2. Two-class Voting

- A variable described as the number of successes in a sequence of independent Bernoulli trials has **Binomial distribution**. Its parameters are n , the number of trials, and p , the probability of success.

- Binomial probability mass function is:

$$P(x) = \mathbf{P}\{X = x\} = \binom{n}{x} p^x q^{n-x}, \quad x = 0, 1, \dots, n,$$

- Binomial distribution:

n	=	number of trials
p	=	probability of success
$P(x)$	=	$\binom{n}{x} p^x q^{n-x}$
$\mathbf{E}(X)$	=	np
$\text{Var}(X)$	=	npq

2. Two-class Voting

- Random Forest: Using 3 Decision Trees

DT1	DT2	DT3	Voting	Probability
0.6	0.6	0.6		
C	C	C	3↑	$P_{\uparrow\uparrow\uparrow}=0.6^3=0.216$
C	C	INC	2↑ 1↓	$P_{\uparrow\uparrow\downarrow}=3\times 0.6^2\times 0.4=0.432$
C	INC	C		
INC	C	C		
INC	INC	C	1↑ 2↓	$P_{\uparrow\downarrow\downarrow}=3\times 0.6\times 0.4^2=0.288$
INC	C	INC		
C	INC	INC		
INC	INC	INC	3↓	$P_{\downarrow\downarrow\downarrow}=0.4^3=0.064$

$$\text{Accuracy} = \frac{0.216+0.432}{1} = 0.648.$$

Using binomial distribution are for any number of votes.

2. Two-class Voting

- Random Forest: Using 51 Decision Trees

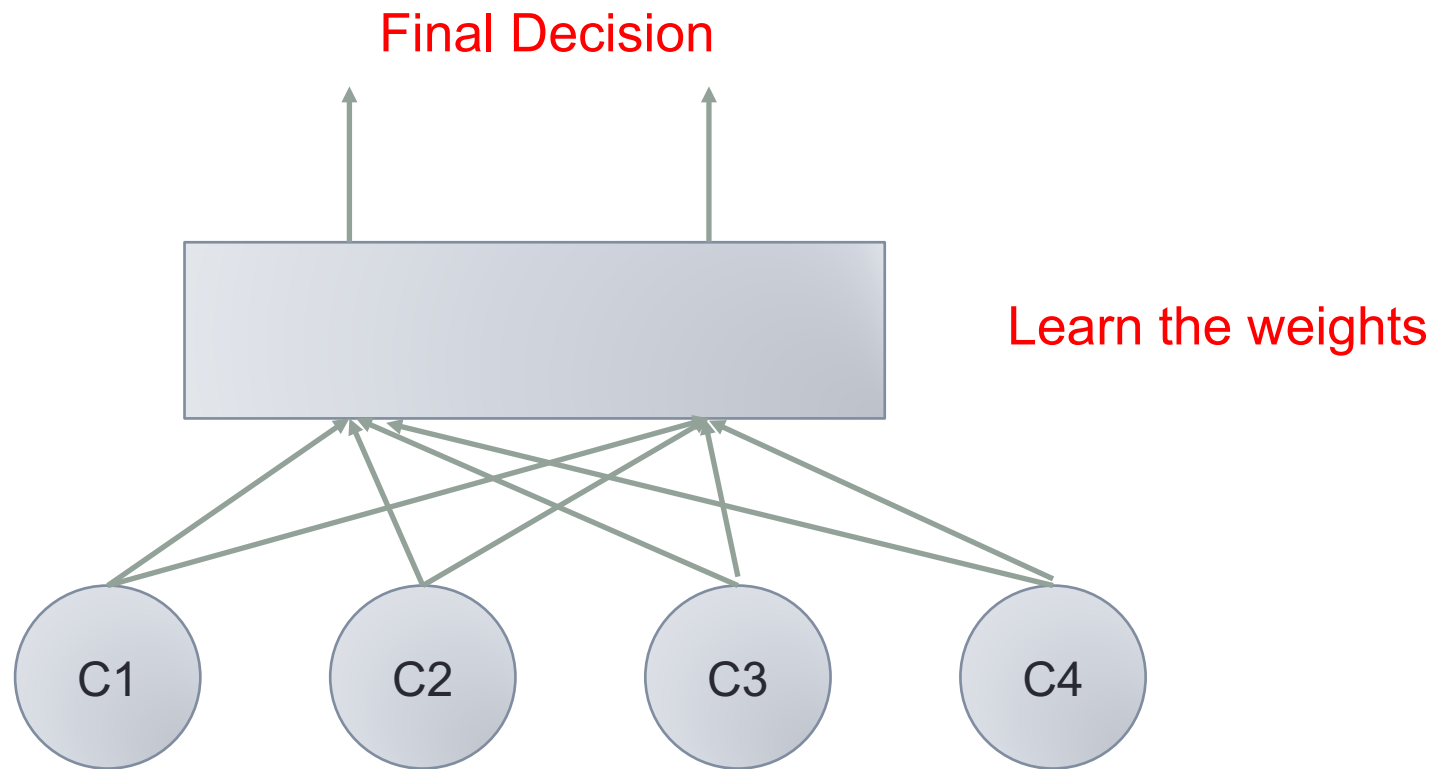
Voting	Probability
51 ↑	$P(51\uparrow) = 0.6^{51} = 4.8497\text{e-}12$
50 ↑ , 1 ↓	$P(50\uparrow, 1\downarrow) = \binom{51}{50} \times 0.6^{50} \times 0.4 = 1.6489\text{e-}10$
.....
26 ↑ , 25 ↓	$P(26\uparrow, 25\downarrow) = \binom{51}{26} \times 0.6^{26} \times 0.4^{25} = 0.0476$
.....
51 ↓	$P(51\downarrow) = 0.4^{51} = 5.0706\text{e-}21$

$$\text{Accuracy} = \frac{P(51\uparrow) + P(50\uparrow, 1\downarrow) + \dots + P(26\uparrow, 25\downarrow)}{1} = 0.9265.$$

Using binomial distribution are for any # of votes.

3. Ensemble Learning

- Example: Voting Machine

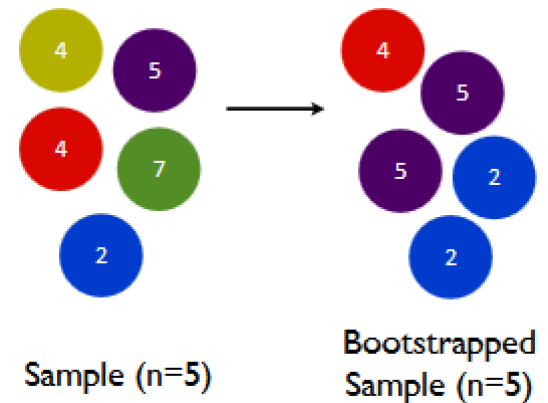

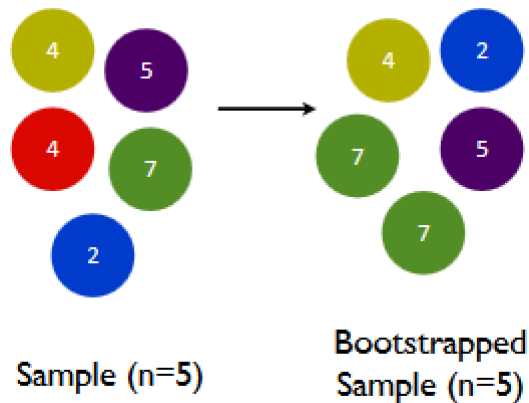


3.1 Bagging

- Bagging = **B**ootstrap **A**ggregating
- In the **Bootstrap**, we replicate our dataset by sampling with replacement:
 - Original Dataset: $Z = (z_1, z_2, \dots, z_N)$, where $z_i = (x_i, y_i)$.
 - Bootstrap samples:
 - $Z^{*1} = \text{sample}(x, 100, \text{replace} = \text{True})$
 -
 - $Z^{*B} = \text{sample}(x, 100, \text{replace} = \text{True})$


3.1 Bagging

- Bootstrap Samples (with replacement):



3.1 Bagging

- Bootstrap Process:



3.1 Bagging

- $S(\mathbf{Z})$ is any quantity computed from the data \mathbf{Z} .
 - For example, its variance:

$$\widehat{\text{Var}}[S(\mathbf{Z})] = \frac{1}{B-1} \sum_{b=1}^B (S(\mathbf{Z}^{*b}) - \bar{S}^*)^2$$

where $\bar{S}^* = \sum_b S(\mathbf{Z}^{*b})/B$

- Apply the bootstrap to estimate prediction error.
 - If $\hat{f}^{*b}(x_i)$ is the predicted value at x_i , from the model fitted to the b -th bootstrap dataset, our estimate is:

$$\widehat{\text{Err}}_{\text{boot}} = \frac{1}{B} \frac{1}{N} \sum_{b=1}^B \sum_{i=1}^N L(y_i, \hat{f}^{*b}(x_i))$$

3.1 Bagging







- For each bootstrap sample Z^{*b} , $b = 1, 2, \dots, B$, we fit our model, giving prediction $\hat{f}^{*b}(x)$.
- The **bagging estimate** is defined by

$$\hat{f}_{\text{bag}}(x) = \frac{1}{B} \sum_{b=1}^B \hat{f}^{*b}(x).$$

- Example: bagging the “linear regression”.
 - Let $\hat{y}^{L,b}$ be the prediction of the decision tree applied to the b -th bootstrap sample.
 - Bagging prediction: $\hat{y}^{\text{boot}} = \frac{1}{B} \sum_{b=1}^B \hat{y}^{L,b}$.
- When a regression method or a classifier has a tendency to **overfit**, Bagging **reduces the variance of the prediction**.


3.1 Bagging

- Example: Tree with simulated data.
 - We generated a sample of size $N = 30$, with two classes and $p = 5$ features.



3.1 Bagging

- Example: Tree with simulated data.



The orange points correspond to the consensus vote, while the green points average the probabilities.

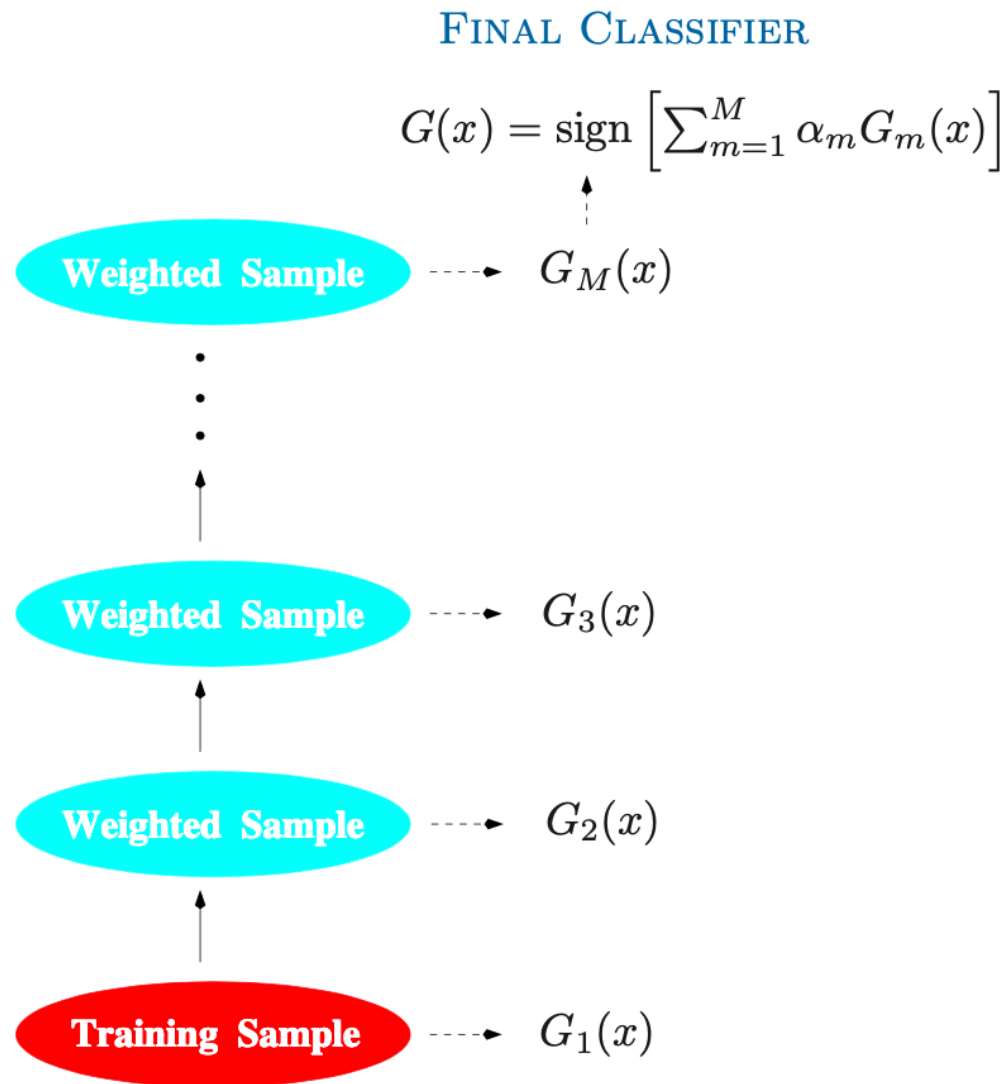
3.2 Boosting

- The motivation for boosting was a procedure that combines the outputs of many “weak” classifiers to produce a powerful “committee.”
- Boosting learns slowly:
 - We first use the samples that are easiest to predict, then slowly down weigh these cases, moving on to harder samples.

3.2 Boosting

- Example: AdaBoost.

The data modifications at each boosting step consist of applying weights w_1, w_2, \dots, w_N to each of the training observations (x_i, y_i) , $i = 1, 2, \dots, N$.



Classifiers are trained on weighted versions of the dataset, and then combined to produce a final prediction.

3.2 Boosting

- Example: Algorithm - AdaBoost.M1

1. Initialize the observation weights $w_i = 1/N$, $i = 1, 2, \dots, N$.

2. For $m = 1$ to M :

(a) Fit a classifier $G_m(x)$ to the training data using weights w_i .

(b) Compute

$$\text{err}_m = \frac{\sum_{i=1}^N w_i I(y_i \neq G_m(x_i))}{\sum_{i=1}^N w_i}.$$

(c) Compute $\alpha_m = \log((1 - \text{err}_m)/\text{err}_m)$.

(d) Set $w_i \leftarrow w_i \cdot \exp[\alpha_m \cdot I(y_i \neq G_m(x_i))]$, $i = 1, 2, \dots, N$.

3. Output $G(x) = \text{sign} \left[\sum_{m=1}^M \alpha_m G_m(x) \right]$.

3.2 Boosting

- Boosted Trees
 - A tree can be formally expressed as

$$T(x; \Theta) = \sum_{j=1}^J \gamma_j I(x \in R_j)$$

- The parameters are found by minimizing the empirical risk

$$\tilde{\Theta} = \arg \min_{\Theta} \sum_{i=1}^N \tilde{L}(y_i, T(x_i, \Theta))$$

3.2 Boosting

- Boosted Trees
 - The boosted tree model is a sum of such trees

$$f_M(x) = \sum_{m=1}^M T(x; \Theta_m)$$

- Where at each step in the **forward stagewise** procedure one must solve

$$\hat{\Theta}_m = \arg \min_{\Theta_m} \sum_{i=1}^N L(y_i, f_{m-1}(x_i) + T(x_i; \Theta_m))$$

3.2 Boosting

- Forward stagewise boosting: **greedy strategy**.
 - At each step the solution tree is the one that maximally reduces the loss, given the current model f_{m-1} and its fits $f_{m-1}(x_i)$.
 - Thus, the tree predictions $T(x_i; \Theta_m)$ are analogous to the components of the **negative gradient**:

$$g_{im} = \left[\frac{\partial L(y_i, f(x_i))}{\partial f(x_i)} \right]_{f(x_i)=f_{m-1}(x_i)}$$

- **Gradient boosting**:
 - Induce a tree $T(x; \Theta_m)$ at the m -th iteration whose predictions t_m are as close **as possible to the negative gradient**.
 - Using **squared error** to measure closeness, this leads us to:

$$\tilde{\Theta}_m = \arg \min_{\Theta} \sum_{i=1}^N (-g_{im} - T(x_i; \Theta))^2$$

3.2 Boosting

- Gradient Tree Boosting Algorithm

1. Initialize $f_0(x) = \arg \min_{\gamma} \sum_{i=1}^N L(y_i, \gamma)$.

2. For $m = 1$ to M :

(a) For $i = 1, 2, \dots, N$ compute

$$r_{im} = - \left[\frac{\partial L(y_i, f(x_i))}{\partial f(x_i)} \right]_{f=f_{m-1}}.$$

(b) Fit a regression tree to the targets r_{im} giving terminal regions R_{jm} , $j = 1, 2, \dots, J_m$.

(c) For $j = 1, 2, \dots, J_m$ compute

$$\gamma_{jm} = \arg \min_{\gamma} \sum_{x_i \in R_{jm}} L(y_i, f_{m-1}(x_i) + \gamma).$$

(d) Update $f_m(x) = f_{m-1}(x) + \sum_{j=1}^{J_m} \gamma_{jm} I(x \in R_{jm})$.

3. Output $\hat{f}(x) = f_M(x)$.

Compute Negative Gradient
with respect to 1,...,m-1 trees

Fit m-th Tree

Update Boosted Tree